


Machine learning regression and marginal effects inference

Rodney Sparapani
Associate Professor of Biostatistics
Medical College of Wisconsin

September 15, 2025

CDC Growth Charts: United States



Motivating Example: Growth Charts

- ▶ US Centers for Disease Control and Prevention (CDC) and the World Health Organization have developed growth charts for childhood development: **height** by **age**, weight by age, body mass index by age and weight by height
- ▶ Here we will focus on **height**, y_t , by **age** in months, $t = 24, \dots, 215$ (2 to 17 years old)
- ▶ CDC uses the LMS method via natural cubic splines (Cole and Green 1992 *Statistics in Medicine*)
- ▶ Three parameters estimated by penalized maximum likelihood the Box-Cox power transformation, L_t ; the mean, M_t ; and the coefficient of variation, S_t

$$z_t = \left\{ \begin{array}{ll} \frac{-1 + (y_t/M_t)^{L_t}}{L_t S_t} & L_t \neq 0 \\ \frac{\log(y_t/M_t)}{S_t} & L_t = 0 \end{array} \right\} \sim N(0, 1)$$

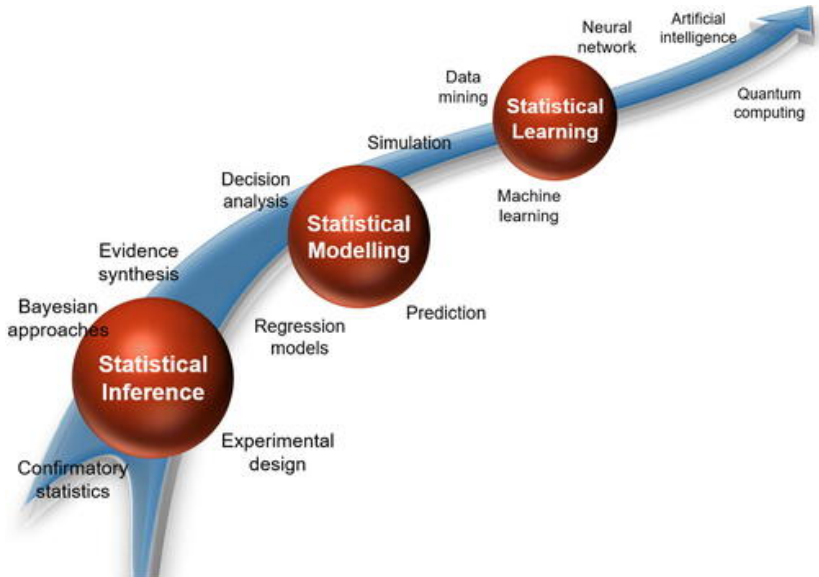
- ▶ But, this only uses part of the data: just males or just females
- ▶ What if we wanted to use all of the data?
- ▶ Or include more information like weight and race/ethnicity?

What is Artificial Intelligence and Statistical Learning?

Artificial intelligence (AI) a computer's ability to perform tasks that normally require human intelligence such as driving a car

- ▶ 1941 (circa): “Machine Intelligence” coined by Alan Turing
- ▶ 1950: Turing’s *Imitation Game* (similar to a *Turing Test*)
- ▶ 1956: “Artificial Intelligence” coined at Dartmouth Workshop
- ▶ 1950 to 2010: AI 1.0, basic research with limited capabilities
- ▶ 2011 to 2017: AI 2.0, deep learning
- ▶ 2018 to today: AI 3.0, foundation/large-language models
- ▶ Howell, Corrado & DeSalvo 2024 *JAMA*

What is Artificial Intelligence and Statistical Learning?



What is Machine Learning (or Statistical Learning)?

- ▶ *Machine learning* is a field within AI
 1. develops methods that learn statistical relationships from *training data* **without being explicitly programmed to do so**
 2. strong learning: a computer's ability to perform its task improves with an increasing pool of *experience*Samuel 1959 (1.) & Mitchell 2006 (2.)
- ▶ For example, you could physically model childhood growth chart data based on principles of human auxology or you could nonparametrically learn the growth curves from training data
- ▶ Back in Samuel's day, linear/logistic regression were considered *machine learning regression (MLR)* for lack of alternatives; however, they do NOT meet the definition due to restrictive linearity and precarious parametric assumptions
- ▶ Linear/logistic regression are proto-MLR rather than MLR
- ▶ Today, by the term "MLR", I mean the widely flexible sense of **without being explicitly programmed to do so** (WBEPTDS)

What are black-box models?

- ▶ The term *black-box*, coined in 1945, for the development of an experimental analysis with electronic circuits that had been in practice about 20 years at that time (Belevitch 1962)
- ▶ Simply ignore the circuit details as-if hidden inside a **black-box** instead, characterize the response output from its stimulus input via experimentation, trial and error, etc.
- ▶ MLR's are typically black-boxes and that is a down-side **a direct interpretation of the model itself is not evident** due to complexity, so let's not bother even trying (in stark contrast to the trivial linear/logistic regression coefficients)
- ▶ In modern terms, a black-box model defies understanding via inspection of the covariates and their associated parameters
- ▶ Rather, an intuitive interpretation is devised by other means such as an orchestrated sequence of covariate setting predictions
- ▶ Therefore, the **rising interest in marginal (*explainable*) effects**
- ▶ Marginal effects are applicable to MLR in general, but here we focus on Bayesian Additive Regression Trees (BART)

What is Machine Learning Regression (MLR)?

- ▶ MLR is extensible, but for the moment consider the general regression case of a continuous outcome with Normal errors

$$y_i = \mu_0 + f(x_i) + \epsilon_i \quad \text{where } \epsilon_i \stackrel{\text{iid}}{\sim} \mathbf{N}(0, \sigma^2)$$

- ▶ f is an unspecified function whose form is to be *learned* from the training data and \mathbf{x}_i is a vector of covariates for $i = 1, \dots, N$ following Samuel's principle via Bayesian nonparametric models without resorting to precarious restrictive assumptions, i.e., we don't want to assume linearity nor pre-specify interactions, etc.

What is Machine Learning Regression (MLR)?

- ▶ *Ensemble learning* discovered in 1997
Krogh & Solich 1997 *Physical Review E*
- ▶ An ensemble of *machines* (in our case binary trees) are fit simultaneously that form the basis of an aggregate prediction with superior performance to any single machine's fit
- ▶ Ensembles are the best currently-known machine learning method with respect to out-of-sample predictive performance for so-called *tabular data* where all of the covariates are of different types, i.e., age, sex, height, weight, etc.
- ▶ N.B. *Deep learning* is inferior to ensembles for tabular data for optimal artificial neural net performance, the inputs need to be all the same type, i.e., all pixels, words or audio waves, etc.
Lundberg and Erion et al. 2020 *Nature Machine Intelligence*
Shwartz-Ziv and Armon 2022. *Information Fusion*

Weighted BART

$$y_i = f(x_i) + \textcolor{blue}{s_i} \epsilon_i \quad \epsilon_i \stackrel{\text{iid}}{\sim} \mathbf{N}(\mathbf{0}, \sigma^2)$$

$$f \stackrel{\text{prior}}{\sim} \text{BART}(\alpha, \beta, H, \kappa, \mu_0, \tau)$$

$$f(x_i) \equiv \mu_0 + \sum_{h=1}^H g(x_i; \mathcal{T}_h, \mathcal{M}_h) \quad H \in \{50, 200, 500\}$$

$$\mu_{hl} | \mathcal{T}_h \stackrel{\text{prior}}{\sim} \mathbf{N}\left(\mathbf{0}, \frac{\tau^2}{4H\kappa^2}\right) \text{ leaves of } \mathcal{T}_h$$
$$\in \mathcal{M}_h$$

$$\sigma^2 \stackrel{\text{prior}}{\sim} \lambda \nu \chi^{-2}(\nu)$$

$$\stackrel{\text{prior}}{\sim} \text{Gamma}^{-1}(\nu/2, \lambda \nu/2) \quad \mathbf{E}[\sigma^2] = \lambda \nu / (\nu - 2)$$

Marginal Effects and Machine Learning Regression (MLR)

- ▶ Suppose we have an MLR, $f(\mathbf{x})$, that is likely a complex function of the covariates with nonlinearities and interactions
- ▶ And we divide the covariates into those of interest, \mathbf{S} , and the complement, \mathbf{C} , not of interest: $f(\mathbf{x}) \equiv f(\mathbf{x}_{\mathbf{S}}, \mathbf{x}_{\mathbf{C}})$
- ▶ Typically, \mathbf{S} is of low-dimension since we intend to peak inside the black-box by visualization: usually 1 to 3 dimensions
- ▶ Let $f_{\mathbf{S}}(\mathbf{x}_{\mathbf{S}})$ denote the marginal effect of $\mathbf{x}_{\mathbf{S}}$

$$\mathbb{E}[y|\mathbf{x}_{\mathbf{S}}] \equiv f_{\mathbf{S}}(\mathbf{x}_{\mathbf{S}})$$

$$f_{\mathbf{S}}(\mathbf{x}_{\mathbf{S}}) \equiv \mathbb{E}_{\mathbf{x}_{\mathbf{C}}} [f(\mathbf{x}_{\mathbf{S}}, \mathbf{x}_{\mathbf{C}}) | \mathbf{x}_{\mathbf{S}}]$$

$$= \int \cdots \int f(\mathbf{x}_{\mathbf{S}}, \mathbf{x}_{\mathbf{C}}) [\mathbf{x}_{\mathbf{C}} | \mathbf{x}_{\mathbf{S}}] d\mathbf{x}_{\mathbf{C}}$$

where $[\mathbf{x}_{\mathbf{C}} | \mathbf{x}_{\mathbf{S}}]$ is the distribution of $\mathbf{x}_{\mathbf{C}} | \mathbf{x}_{\mathbf{S}}$

$$\approx \int \cdots \int f(\mathbf{x}_{\mathbf{S}}, \mathbf{x}_{\mathbf{C}}) [\mathbf{x}_{\mathbf{C}}] d\mathbf{x}_{\mathbf{C}} \quad \text{assuming } \mathbf{x}_{\mathbf{S}} \perp \mathbf{x}_{\mathbf{C}}$$

Friedman's partial dependence function (FPD) and Marginal Effects Assuming Independent Covariates

$$\mathbb{E}[y|\mathbf{x}_S] \equiv f_S(\mathbf{x}_S)$$

$$f_S(\mathbf{x}_S) \equiv \mathbb{E}_{x_C} [f(\mathbf{x}_S, x_C) | \mathbf{x}_S]$$

$$\approx N^{-1} \sum_i f(\mathbf{x}_S, x_{iC}) \quad \text{the partial dependence function}$$

where x_{iC} are the training values

$$f_{Sm}(\mathbf{x}_S) \approx N^{-1} \sum_i f_m(\mathbf{x}_S, x_{iC}) \quad \text{for each MCMC sample } m$$

$$\hat{f}_S(\mathbf{x}_S) \approx M^{-1} \sum_m f_{Sm}(\mathbf{x}_S) \quad \text{estimate by averaging over the posterior}$$

Friedman 2001 *Annals of Statistics*


Friedman's partial dependence function (FPD) and Marginal Effects Assuming Independent Covariates

Linear regression example: age = $x_S = x_1$, weight = $x_C = x_2$


$$\begin{aligned}y_i &= \beta_0 + \beta_1 x_{1i} + \beta_2 x_{2i} + \epsilon_i \\f(x_{1i}, x_{2i}) &= \beta_0 + \beta_1 x_{1i} + \beta_2 x_{2i} \\f_S(x_1) &= E_{x_2} [f(x_1, x_{2i}) | x_1] \\&= E_{x_2} [\beta_0 + \beta_1 x_1 + \beta_2 x_{2i} | x_1] \\&= \beta_0 + \beta_1 x_1 + \beta_2 E_{x_2} [x_{2i}] \\&\approx N^{-1} \sum_i f(x_1, x_{2i}) \\&= N^{-1} \sum_i (\beta_0 + \beta_1 x_1 + \beta_2 x_{2i}) \\&= \beta_0 + \beta_1 x_1 + \beta_2 \bar{x}_2\end{aligned}$$

Friedman's partial dependence function (FPD) and Marginal Effects Assuming Independent Covariates

Linear regression example: age = $x_S = x_1$, weight = $x_C = x_2$



Weighted BART: s_i for height and \tilde{s}_i for weight



Independent Covariates:

Simple Random Sampling (SRS) marginal

Lundberg & Lee 2017 *NIPS*

Janzing et al. 2020 *PMLR*

- ▶ To speed up FPD when N is large
- ▶ Consider our growth chart for height example
age = t , sex = u , race = v and weight = w
 $y_i = f(t_i, u_i, v_i, w_i) + s_i \epsilon_i$ where $f \stackrel{\text{prior}}{\sim} \text{BART}$
- ▶ $S = (t, u)$ and $C = (v, w)$
- ▶ Like FPD, but with K random draws of (v_i, w_i) from training
- ▶ K does not have to be large, e.g., $K = 30$ is often sufficient

$$\begin{aligned} f_{t,u}(t, u) &= \mathbb{E}_{v,w} [f(t, u, v, w) | t, u] \\ &\approx K^{-1} \sum_k f(t, u, v_k^*, w_k^*) \end{aligned}$$

where $(v_k^*, w_k^*) \sim \{(v_1, w_1), \dots, (v_N, w_N)\}$ simple random sample

Dependent Covariates: the Monte Carlo (MC) marginal

- ▶ Consider our growth chart for height example
age = t , sex = u , race = v and weight = w
 $y_i = f(t_i, u_i, v_i, w_i) + s_i \epsilon_i$ where $f \stackrel{\text{prior}}{\sim} \text{BART}$
- ▶ Next we model the strong relationship between covariates
 $E[w|t, u, v] = \tilde{w} = \tilde{f}(t, u, v)$
- ▶ Here we summarize this relationship with a BART model
 $w_i = \tilde{f}(t_i, u_i, v_i) + \tilde{s}_i \tilde{\epsilon}_i$ where $\tilde{f} \stackrel{\text{prior}}{\sim} \text{BART}$
- ▶ For marginal effects applicable to dependent variables
- ▶ $S = (t, u)$ and $C = (v, w)$

$$\begin{aligned} f_{t,u}(t, u) &= E_{v,w} [f(t, u, v, w) | t, u] \\ &\approx K^{-1} \sum_k f(t, u, v_k^*, w_k^*) \end{aligned}$$

where $v_k^* \sim \{v_1, \dots, v_N\}$

simple random sample

and $w_k^* \sim N(\tilde{f}(t, u, v_k^*), \sigma_{\tilde{\epsilon}}^2)$

Dependent Covariates: a hybrid marginal

- ▶ Consider our growth chart for height example
age = t , sex = u , race = v and weight = w
 $y_i = f(t_i, u_i, v_i, w_i) + s_i \epsilon_i$ where $f \stackrel{\text{prior}}{\sim} \text{BART}$
- ▶ Next we model the strong relationship between covariates
 $E[w|t, u] = \tilde{w} = \tilde{f}(t, u)$
- ▶ Here we summarize this relationship with a BART model
 $w_i = \tilde{f}(t_i, u_i) + \tilde{s}_i \tilde{\epsilon}_i$ where $\tilde{f} \stackrel{\text{prior}}{\sim} \text{BART}$
- ▶ For marginal effects applicable to dependent variables
- ▶ $S = (t, u)$ and $C = (v, w)$

$$\begin{aligned} f_{t,u}(t, u) &= E_{v,w} [f(t, u, v, w) | t, u] \\ &\approx K^{-1} \sum_k f(t, u, v_k^*, w_k^*) \end{aligned}$$

where $v_k^* \sim \{v_1, \dots, v_N\}$


and $w_k^* \sim N(\tilde{f}(t, u), \sigma_{\tilde{\epsilon}}^2)$

simple random sample

Dependent Covariates: the Synthetic Approximation (SA) marginal

Linear regression example: age = $x_S = x_1$, weight = $x_C = x_2$

$$f_S(x_S) \approx f(x_S, \mathbf{E}_{x_C}[x_C|x_S])$$



Dependent Covariates: the Synthetic Approximation (SA) marginal

- ▶ We proceed as before by modeling the strong relationship between age, sex and weight first
- ▶ For marginal effects applicable to dependent variables
- ▶ $S = (t, u)$ and $C = (v, w)$

$$E[w|t, u] = \tilde{f}(t, u)$$

$$w_i = \tilde{f}(t_i, u_i) + \tilde{s}_i \tilde{\epsilon}_i \text{ where } \tilde{f} \stackrel{\text{prior}}{\sim} \text{BART}$$

$$\widehat{w}(t, u) = M^{-1} \sum_m \tilde{f}_m(t, u)$$

$$\begin{aligned} f_{t,u}(t, u) &= E_v[f(t, u, v, w)|t, u, w = E[w|t, u]] \\ &= E_v[f(t, u, v, \widehat{w}(t, u))|t, u] \\ &\approx N^{-1} \sum_i f(t, u, v_i, \widehat{w}(t, u)) \end{aligned}$$

Dependent Covariates: the Nearest Neighbor (NN) marginal

- ▶ t for age, u for sex, v for race/ethnicity and w for weight
- ▶ For age, t , we have chosen this grid of values
 $-\infty = t_0^* < t_1^* < t_2^* < \dots < t_J^* < t_{J+1}^* = \infty$
where $t_1^* = 2, \dots, t_{17}^* = 18$
- ▶ For sex, u , we have just two values: $u^* \in \{M, F\}$

$$f_{\textcolor{red}{S}}(\textcolor{red}{t}_j^*, u^*) \approx K(t_j^*, u^*)^{-1} \sum_{\mathcal{X}(t_j^*, u^*)} f(\textcolor{red}{t}_j^*, u^*, v_i, \textcolor{blue}{w}_i)$$

where $\mathcal{X}(t_j^*, u^*) = \{i : t_{j-1}^* < t_i < t_{j+1}^*, u_i = u^*\}$


and $K(t_j^*, u^*) = |\mathcal{X}(t_j^*, u^*)|$

Returning to the real data example

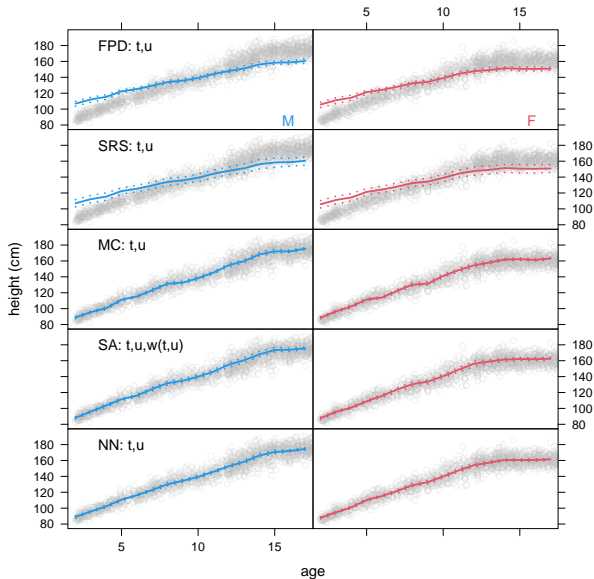
- ▶ CDC's data is the US National Health and Nutrition Examination Survey (NHANES) waves I-III circa 1972 (I), 1978 (II), 1991 (III): $n=12677$
- ▶ For simplicity, I used NHANES annual/continuous 1999-2000
- ▶ The data set is in the BART3 package: `bm` see the `growth*.R` examples in demo
- ▶ 2-17 years (fractional age for months)
- ▶ each child only measured once
- ▶ height (cm) and weight (kg) collected

	<i>n</i>	%
Total	3435	
Males	1768	51.5
Females	1667	48.5
White	800	23.3
Black	1035	30.1
Hispanic	1600	46.6

$R^2 = 96.2\%$ in the testing subset: growth1.R



Marginal effect of age comparison



MLR marginal effects and computational efficiency

- ▶ *Shapley values* are a popular choice for explainability research that are based on marginal effects
- ▶ However, **Shapley values are very computationally intensive**
In terms of complexity, they are considered to be NP-hard: not practical unless the number of covariates is small
- ▶ Shapley values approximate $f(\mathbf{x})$ by additive effects (typically of one variable at a time), e.g., $f(\mathbf{x}) \approx \mu_0 + \sum_j f_j(x_j)$
- ▶ *$f(\mathbf{x})$ is additive in terms of single covariate functions, $f_j(x_j)$, i.e., effectively, we are assuming independence*
- ▶ But there is a common extension for two-way interactions
Lundberg and Erion et al. 2020 *Nature Machine Intelligence*
- ▶ And multi-way interactions have been defined as well
Grabisch & Roubens 1999; Borgonovo et al. 2024

Shapley value marginal effects of Independent Covariates

- ▶ **Two equivalent definitions: original ordered vs. more computationally friendly unordered**
- ▶ $S = \{x_j\}$ and C_{-j} contains all $P - 1$ other covariates
- ▶ \mathcal{P}_j is the set of all *ordered* permutations of $S \cup C_{-j}$
 $f_j(x_j) \equiv (P!)^{-1} \sum_{S_* \in \mathcal{P}_j} [f_{+j}^*(x_{S_*}) - f_{-j}^*(x_{S_*})]$
where $f_{+j}^*(x_{S_*})$ evaluates arguments up to/including x_j
while $f_{-j}^*(x_{S_*})$ evaluates arguments before/excluding x_j
- ▶ C_j is the set of all *unordered* combinations $S_* \subset C_{-j}$
 $f_j(x_j) \equiv \sum_{S_* \in C_j} \frac{|S_*|!(P-|S_*|-1)!}{P!} [f_S(x_{S_*}, x_j) - f_S(x_{S_*})]$
- ▶ If each $f_S(\cdot)$ are fit from the training
the number of fits needed grows rapidly with P

P	2	3	4	5	10	20	30	P
Fits	3	7	15	31	1,023	1,048,575	1,073,741,823	$2^P - 1$

Fast Shapley value approximations from a single fit

- ▶ Rather than fitting so many models, Shapley values can be created from a single fit's marginal effects
- ▶ For example, suppose $f_S(x_S) = E_{x_C} [f(\mathbf{x}_S, \mathbf{x}_C) | \mathbf{x}_S]$
- ▶ This would certainly help but the computations are still daunting unless the number of covariates is small
- ▶ There is a simple algorithm, known as EXPVALUE, for these marginals that is basically equivalent to FPD
And there are more efficient, so-called Tree SHAP, algorithms but these are far more complex to program
Lundberg and Erion et al. 2020 *Nature Machine Intelligence*
- ▶ And advanced random sampling schemes have been proposed but they are challenging to implement as well
Yang, Zhou et al. 2023 *JASA*

Shapley value marginals for **Dependent Covariates**

Marginal effect of age

- ▶ Shapley values come from game theory where each player takes their turn and the order of play is important
- ▶ The *players* here are the covariates
- ▶ And as can be shown, the order of covariates doesn't really matter i.e., the order of covariates is arbitrary (Lundberg and Lee 2017)
- ▶ Nevertheless, all possible orderings of t, u, v, w : $P! = 24$

age first	age second	age third	age last
t, u, v, w	u, t, v, w	u, v, t, w	u, v, w, t
t, u, w, v	u, t, w, v	u, w, t, v	u, w, v, t
t, v, u, w	v, t, u, w	v, u, t, w	v, u, w, t
t, v, w, u	v, t, w, u	v, w, t, u	v, w, u, t
t, w, u, v	w, t, u, v	w, u, t, v	w, u, v, t
t, w, v, u	w, t, v, u	w, v, t, u	w, v, u, t

Shapley value marginal effects of Dependent Covariates

Marginal effect of age

Differentials for t corresponding to each ordering

$$f(t)-0 \quad f(u,t)-f(u) \quad f(u,v,t)-f(u,v) \quad f(u,v,w,t)-f(u,v,w)$$

$$f(t)-0 \quad f(u,t)-f(u) \quad f(u,w,t)-f(u,w) \quad f(u,w,v,t)-f(u,w,v)$$

$$f(t)-0 \quad f(v,t)-f(v) \quad f(v,u,t)-f(v,u) \quad f(v,u,w,t)-f(v,u,w)$$

$$f(t)-0 \quad f(v,t)-f(v) \quad f(v,w,t)-f(v,w) \quad f(v,w,u,t)-f(v,w,u)$$

$$f(t)-0 \quad f(w,t)-f(w) \quad f(w,u,t)-f(w,u) \quad f(w,u,v,t)-f(w,u,v)$$

$$f(t)-0 \quad f(w,t)-f(w) \quad f(w,v,t)-f(w,v) \quad f(w,v,u,t)-f(w,v,u)$$

Weighted differentials for t corresponding to each ordering

$$6f(t) \quad 2[f(t,u)-f(u)] \quad 2[f(t,u,v)-f(u,v)] \quad 6[f(t,u,v,w)-f(u,v,w)]$$

$$2[f(t,v)-f(v)] \quad 2[f(t,u,w)-f(u,w)]$$

$$2[f(t,w)-f(w)] \quad 2[f(t,v,w)-f(v,w)]$$

0	1	2	3
3!	2!	2!	3!

Last row are the weights for the differentials: $|S_*|!(P - |S_*| - 1)!$

(Lundberg and Lee 2017)

Shapley values and

Marginal Effects for Dependent Covariates:

Synthetic Approximation marginal

- As before, rely on the strong relationships of age, sex and weight


$$E[w|t, u] = \tilde{f}(t, u)$$

$$w_i = \tilde{f}(t_i, u_i) + \tilde{\epsilon}_i \text{ where } \tilde{f}^{\text{prior}} \sim \text{BART}$$

- For a marginal effect more applicable to dependent variables

$$\text{Females: } \mu + f_t(t) + f_u(\mathbf{F}) + 2f_{t:u}(t, \mathbf{F}) + f_w(\widehat{w}(t, \mathbf{F}))$$

Marginal effects with Shapley values



Marginal effect of age: computational efficiency measured by `system.time()` in seconds

Method	Computational Timings			
	user		elapsed	
	s	%	s	%
Synthetic approximation	340	100	64	100
Nearest neighbor	32	9	20	31
SRS: $K = 30$	130	38	17	27
SRS: $K = 5$	22	6	3	5
SV: t , age-only	1610		1610	
SV: u , sex-only	249		249	
SV: w , weight-only	2007		2011	
SV: Synthetic approximation	3866	1137	3870	6047

Marginal effects for dependent covariates and computational efficiency

- ▶ At first, it is quite surprising that FPD assumes independence since it has the term *dependence* in its name
- ▶ Our novel marginals Nearest Neighbor and Synthetic Approximation are computationally efficient
- ▶ Yet the Shapley value marginals are very computationally demanding and often impractical
- ▶ It is possible to exploit the structure of binary trees to compute Shapley values by the so-called Tree SHAP algorithms Lundberg and Erion et al. 2020 *Nature Machine Intelligence*
- ▶ For example, see the **treeshap** R package for Random Forests but this still might not be fast enough to be feasible
- ▶ Furthermore, this has not been adapted to BART FWIW

Marginal effects for dependent covariates and computational efficiency

- My **BART3** package on github has S3 methods for marginals

S3 method	Assumes \perp	Marginal type
FPD	Yes	FPD
SRS	Yes	Simple random sampling
SHAP	Yes	SV
SHAP2	Yes	SV two-way interaction
NN	No	Nearest neighbors
SHNN	No	SV with nearest neighbors
SHNN2	No	SV two-way interaction with nearest neighbors

Conclusion

- ▶ This was an overview of BART and its place in machine learning
- ▶ Our focus was on the BART prior for continuous outcomes
- ▶ In particular, estimating marginal effects with BART whether assuming independence or dependence
- ▶ We contrasted Friedman's partial dependence function with Shapley values
- ▶ And we have described facilitating these calculations with opportunities for bettering performance statistically and computationally
- ▶ We provide a reference implementation in the **BART3** R package with *new and improved* marginal effects S3 functions