

Monotonic BART and outlier detection

Rodney Sparapani
Associate Professor of Biostatistics
Medical College of Wisconsin


September 16, 2025

Chronically ill children and potential height outliers

- ▶ This data is from the electronic health records (EHR) of a large regional children's health care system
- ▶ Chronically ill children are often at high risk for malnutrition
- ▶ Typically this is assessed by comparison to US Centers for Disease Control (CDC) growth chart benchmarks (or WHO)
- ▶ CDC inputs are age, gender, height and weight
- ▶ Age and gender are extremely reliable
- ▶ However, height and weight are prone to EHR outliers and there is practically NO quality control for these measures i.e., **the ground truth of height outliers is largely unknown**
- ▶ Based on a UK study, there are an order of magnitude more height (3%) than weight outliers (0.2%) **per measurement** (Phan et al. 2020 Scientific Reports)
- ▶ **Determining malnourishment requires height outlier detection**
- ▶ Furthermore, this method should be robust to weight outliers that are harder to identify (yet fortunately less prevalent)

<https://www.cdc.gov/growthcharts>

CDC Growth Charts: United States



Monotonic BART (mBART)





Chipman et al. 2021 *Bayesian Analysis*

- ▶ $f \stackrel{\text{prior}}{\sim} \text{mBART}$
- ▶ A function f is monotone with respect to x_j if f satisfies
$$f(\dots, x_{j-1}, x_j + \Delta x, x_{j+1}, \dots) \geq f(\dots, x_{j-1}, x_j, x_{j+1}, \dots)$$
for all $\Delta x > 0$ (increasing/nondecreasing) or
for all $\Delta x < 0$ (decreasing/nonincreasing)
- ▶ Constraint Conditions for Tree Monotonicity
A tree function $g(\mathbf{x}; \mathcal{T}, \mathcal{M})$ will be monotone in coordinate x_j if the leaf value of each of its terminal node regions is
 - not greater than the minimum level of all of its above-neighbor regions with respect to x_j and
 - not less than the maximum level of all of its below-neighbor regions with respect to x_j

Monotonic BART (mBART)

- ▶ dichotomous covariates are monotonic in the trivial sense
- ▶ for continuous outcomes, we can restrict to monotonicity
- ▶ e.g, assume age, t , is monotonic with respect to height, y
- ▶ do we want to assume the same for age, t , and weight, w ?
- ▶ we can relax this restriction like so: $x = (t, u, v, w, -w)$
- ▶ yet from here on, we assume that weight is monotonic

Monotonic example: increasing in x_1 and x_2



Monotonic BART (mBART)

Chipman et al. 2021 *Bayesian Analysis*

- ▶ The leaf prior for BART $\mu_j | \mathcal{T}^{\text{prior}} \sim N(0, \sigma_\mu^2)$
- ▶ Consider the simplest case of two monotonic leaves in mBART (relying on the results of Azzalini 1985 *Scand J Stat*)

$$\begin{bmatrix} \mu_1 \\ \mu_2 \end{bmatrix}^{\text{prior}} \sim N_2 \left(\vec{0}_2, c^2 \sigma_\mu^2 I_2 \right) \mathbf{I}(\mu_1 < \mu_2) \text{ where } c^2 = \frac{\pi}{\pi - 1} \approx 1.47$$

Equivalent to skew Normal marginals where $V[\mu_1] = V[\mu_2] = \sigma_\mu^2$

$$\begin{aligned} \mu_1^{\text{prior}} &\sim \phi\left(\frac{\mu_1}{c\sigma_\mu}\right) \Phi\left(\frac{-\mu_1}{c\sigma_\mu}\right) & E[\mu_1] &= \frac{-\sigma_\mu}{\sqrt{\pi - 1}} \\ \mu_2^{\text{prior}} &\sim \phi\left(\frac{\mu_2}{c\sigma_\mu}\right) \Phi\left(\frac{\mu_2}{c\sigma_\mu}\right) & E[\mu_2] &= \frac{+\sigma_\mu}{\sqrt{\pi - 1}} \end{aligned}$$

Nonparametric outlier detection

- ▶ *Monotonicity provides additional robustness to outliers since f can't just rise towards an outlier and double-back down afterward (nor in the inverse directions)*
- ▶ We have *population* predictions of the form $\hat{y}_{ij} = E[y_{ij}] = \hat{f}(x_{ij})$ where $j = 1, \dots, n_i$
- ▶ *But these expectations are biased except for the average child*
- ▶ We need to adjust these up or down for a given subject
- ▶ So let $m_i = \text{median}_j(y_{ij} - \hat{y}_{ij})$
(median rather than mean to be robust to outliers)
- ▶ Now, we make *personalized* predictions $\tilde{y}_{ij} = m_i + \hat{y}_{ij}$
- ▶ We define the *relative error* of these as $d_{ij} = (y_{ij} - \tilde{y}_{ij})/\tilde{y}_{ij}$
- ▶ Outliers are defined as $|d_{ij}| > \delta$ where δ can be determined from the Receiver Operating Characteristic (ROC) curve
- ▶ And the discriminating performance of the method is assessed by the area under the ROC curve

Friedman's partial dependence function and the Synthetic Approximation marginal

- ▶ t for age, u for sex, v for race and w for weight
- ▶ To do this right, first consider the likely monotonic relationship between age, sex and weight
 $E[w|t, u] = \tilde{f}(t, u)$
- ▶ So we need to fit an intermediate mBART model first
 $w_i = \tilde{f}(t_i, u_i) + \tilde{\epsilon}_i$ where $\tilde{f} \stackrel{\text{prior}}{\sim} \text{mBART}$
- ▶ The Synthetic Approximation marginal for dependent variables

$$\begin{aligned} f_{t,u}(t, u) &= E_v [f(t, u, v, w^*) | t, u, w^* = E[w|t, u]] \\ &= E_v [f(t, u, v, \tilde{f}(t, u)) | t, u] \end{aligned}$$

Returning to the real data example

Sparapani et al. 2022 *JPGN*

- ▶ Constructed two independent cohorts of chronically ill children
 - ▶ 2-8 years old
 - ▶ measured at least every 120 days on average
 - ▶ followed for at least 2 years
- ▶ Training cohort: 1376 children with height outliers unknown
39491 measurements: 28.7/child on average
- ▶ Validation cohort: 318 children
7378 measurements: 23.2/child on average
manually reviewed to determine height outliers
however, the *ground truth* is fallible
it is retrospective: we can't just re-measure these children
- ▶ Heights in the training cohort fit with mBART to
age, sex, race/ethnicity and weight


Returning to the real data example

- ▶ Outlier detection conducted for the Validation cohort
- ▶ The area under the Receiver Operating Characteristic (ROC) curve was **excellent: 0.841**
- ▶ By comparison, if you use the CDC height by age growth chart, the area is only 0.776
- ▶ Based on ROC curve, two relative error cutoffs considered **Aggressive, 0.075**; and Conservative, 0.085


Real data summary

	Training 1376	Validation 318
Children	<i>n</i> (%)	<i>n</i> (%)
Female	594 (43.2%)	132 (41.5%)
Black	313 (22.7%)	66 (20.8%)
White	783 (56.9%)	189 (59.4%)
Other	280 (20.3%)	63 (19.8%)
Children with outliers	N/A	101 (31.8%)
Measurements	<i>m</i>	<i>m</i>
Height (cm)	39491	7378
	Mean (SD)	Mean (SD)
Measurements/child	28.7	23.2
First visit age 2	86.4 (8.8)	84.5 (6.6)
Last visit age 5	111.3 (8.4)	109.5 (9.4)
	%	%
mBART R^2	82.2%	75.3%


FPD marginal of age: assuming Independence




SA marginal of age: assuming Dependence




FPD marginal of weight: assuming Independence



SA marginal of weight: assuming **Dependence**



Receiver Operating Characteristic (ROC) curve



Aggressive cutoff 0.075

- ▶ B: mBART outlier detection
- ▶ C: clinical review ground truth
- ▶ Outlier: 0 (False), 1 (True)

	B=0	B=1	
C=0	TN=172	FP=45	M=217
C=1	FN=29	TP=72	Q=101
	N=201	P=117	T=318

$$\text{Sensitivity or Recall} = \mathbf{P[B = 1|C = 1]} = \frac{TP}{Q} = \frac{72}{101} = 0.713$$

$$\text{Specificity} = \mathbf{P[B = 0|C = 0]} = \frac{TN}{M} = \frac{172}{217} = 0.793$$

$$\text{PPV or Precision} = \mathbf{P[C = 1|B = 1]} = \frac{TP}{P} = \frac{72}{117} = 0.615$$

$$\text{NPV} = \mathbf{P[C = 0|B = 0]} = \frac{TN}{N} = \frac{172}{201} = 0.856$$

Conservative cutoff 0.085

- ▶ B: mBART outlier detection
- ▶ C: clinical review ground truth
- ▶ Outlier: 0 (False), 1 (True)

	B=0	B=1	
C=0	TN=186	FP=31	M=217
C=1	FN=37	TP=64	Q=101
	N=223	P=95	T=318

$$\text{Sensitivity or Recall} = \mathbf{P[B = 1|C = 1]} = \frac{TP}{Q} = \frac{64}{101} = \mathbf{0.634}$$

$$\text{Specificity} = \mathbf{P[B = 0|C = 0]} = \frac{TN}{M} = \frac{186}{217} = \mathbf{0.857}$$

$$\text{PPV or Precision} = \mathbf{P[C = 1|B = 1]} = \frac{TP}{P} = \frac{64}{95} = \mathbf{0.674}$$

$$\text{NPV} = \mathbf{P[C = 0|B = 0]} = \frac{TN}{N} = \frac{186}{223} = \mathbf{0.834}$$

Aggressive cutoff: targeted smoothing BART with monotonic weight

Starling et al. Annals of Applied Statistics 2020

- ▶ B: mBART outlier detection
- ▶ C: clinical review ground truth
- ▶ Outlier: 0 (False), 1 (True)

	B=0	B=1	
C=0	TN=165	FP=52	M=217
C=1	FN=27	TP=74	Q=101
	N=192	P=126	T=318

$$\text{Sensitivity or Recall} = P[B = 1|C = 1] = \frac{TP}{Q} = \frac{74}{101} = 0.732$$

$$\text{Specificity} = P[B = 0|C = 0] = \frac{TN}{M} = \frac{165}{217} = 0.760$$

$$\text{PPV or Precision} = P[C = 1|B = 1] = \frac{TP}{P} = \frac{74}{126} = 0.587$$

$$\text{NPV} = P[C = 0|B = 0] = \frac{TN}{N} = \frac{165}{192} = 0.859$$

Aggressive cutoff: females only

- ▶ B: mBART outlier detection
- ▶ C: clinical review ground truth
- ▶ Outlier: 0 (False), 1 (True)

	B=0	B=1	
C=0	TN=70	FP=20	M=90
C=1	FN=11	TP=31	Q=42
	N=81	P=51	T=132

$$\text{Sensitivity or Recall} = P[B = 1|C = 1] = \frac{TP}{Q} = \frac{31}{42} = 0.738$$

$$\text{Specificity} = P[B = 0|C = 0] = \frac{TN}{M} = \frac{70}{90} = 0.778$$

$$\text{PPV or Precision} = P[C = 1|B = 1] = \frac{TP}{P} = \frac{31}{51} = 0.608$$

$$\text{NPV} = P[C = 0|B = 0] = \frac{TN}{N} = \frac{70}{81} = 0.864$$

Aggressive cutoff: non-whites only

- ▶ B: mBART outlier detection
- ▶ C: clinical review ground truth
- ▶ Outlier: 0 (False), 1 (True)

	B=0	B=1	
C=0	TN=60	FP=19	M=79
C=1	FN=11	TP=39	Q=50
	N=71	P=58	T=129





$$\text{Sensitivity or Recall} = P[B = 1|C = 1] = \frac{TP}{Q} = \frac{39}{50} = 0.780$$

$$\text{Specificity} = P[B = 0|C = 0] = \frac{TN}{M} = \frac{60}{79} = 0.759$$


$$\text{PPV or Precision} = P[C = 1|B = 1] = \frac{TP}{P} = \frac{39}{58} = 0.672$$

$$\text{NPV} = P[C = 0|B = 0] = \frac{TN}{N} = \frac{60}{71} = 0.845$$


True Positives




False Positives




(a) Sex M, Race W



(b) Sex M, Race W




(c) Sex M, Race B




(d) Sex M, Race B


False Negatives




(a) Sex F, Race W



(b) Sex F, Race W



(c) Sex M, Race W







(d) Sex M, Race W





Conclusions

- ▶ We constructed our new outlier detection methodology based on nonparametric machine learning with monotonic BART
- ▶ This automated method's performance was deemed to be adequate via an independent validation cohort
- ▶ Modern methodology leads to a simply-tuned single rule as opposed to complex simultaneous tuning of multiple rules that have been proposed (by others) based on classic methods
- ▶ For EHR heights/weights, the ground truth is unknown prospective corrections are rarely performed and retrospective attempts to identify outliers manually are fallible
- ▶ For covariates that have dependent relationships, we provide an extension to compute the marginal effects for one, or more, of these variables at a time: this extension has wide applicability in nonparametric/machine learning regression

True Positives



False Positives



False Negatives

